



# CSAI 801 Project: COVID-19 Outcome Prediction

Mohamed Abdelrazek Rezaik Salim

ID: 20399127

Supervised by:

Dr. Marwa Elsayed

Eng. Amr Zaki

## Project Description

- The data used in this project will help to identify whether a person is going to recover from coronavirus symptoms or not based on some pre-defined standard symptoms. These symptoms are based on guidelines given by the World Health Organization (WHO).
- This dataset has daily level information on the number of affected cases, deaths and recovery from 2019 novel coronavirus. Please note that this is a time series data and so the number of cases on any given day is the cumulative number.
- The data is available from 22 Jan, 2020. Data is in “data.csv”.

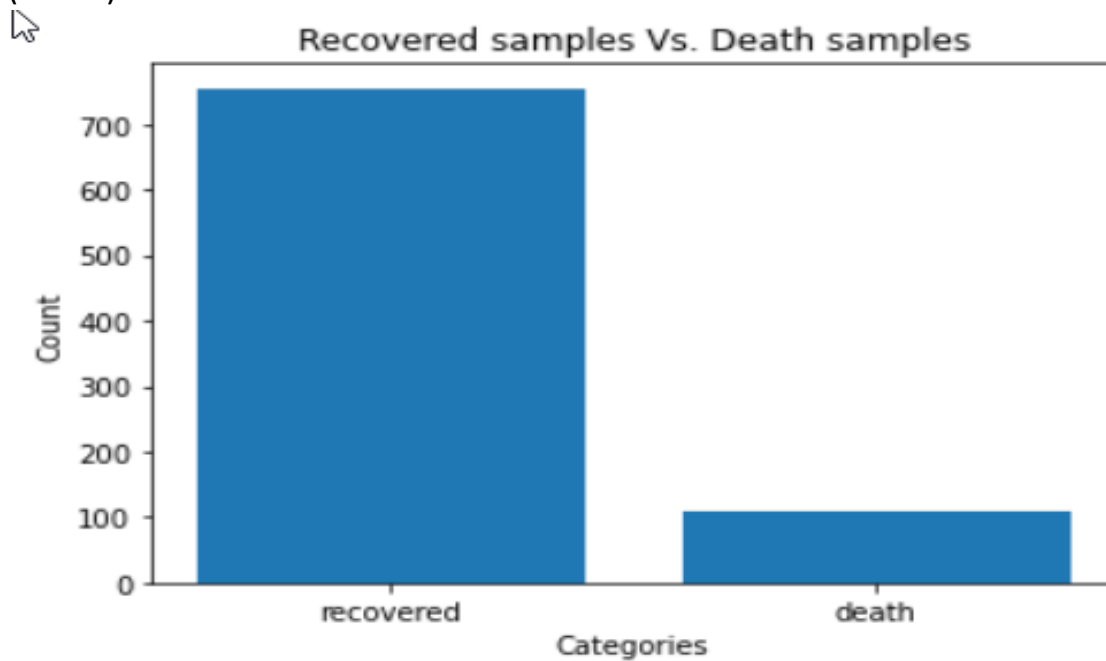
## Dataset Properties

The dataset contains 14 major variables that will be having an impact on whether someone has recovered or not, the description of each variable are as follows,


- 1) **Country**: where the person resides
- 2) **Location**: which part in the Country
- 3) **Age**: Classification of the age group for each person, based on WHO Age Group Standard
- 4) **Gender**: Male or Female
- 5) **Visited Wuhan**: whether the person has visited Wuhan, China or not
- 6) **From Wuhan**: whether the person is from Wuhan, China or not
- 7) **Symptoms**: there are six families of symptoms that are coded in six fields.
- 8) **Time before symptoms appear**:
- 9) **Result**: death (1) or recovered (0)

## Exploratory Data Analysis (EDA)


- The data is already cleaned as it is mentioned in the description.
- The data has 13 features and one target column.
- The data has no null values, no outliers, and columns data types are correct.
- The data is already ordinal encoded.
- The data has 863 Records.
- The data is imbalanced, where the percentage of class 1 (recovered) is 13%, and class 0 (death) is 87%.



- The age feature shows that the range from 40 to 60 of age represents most patients.



- Some features have a strong correlation with the target column, and others do not.



- As we can see the “age” column has the strongest relationship with the target it’s about **+0.5** and “diff\_sym\_hos” is about **+0.22**.
- So, we can assume the age column has a big effect on the target column.

## Splitting and Preparing the Data for Training

- We will split the data into training and testing 0.8 to 0.2 respectively with a stratified distribution.

```
# get samples and target
X = df[features]
y = df["result"]
# Get train and test data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, np.ravel(y),
                                                    test_size=0.2,
                                                    shuffle=True,
                                                    random_state=1,
                                                    stratify=y)
```

- Because the data is imbalanced, we will use the oversampling technique to reduce the imbalance effect, but we will not use it with all types of algorithms.

```
from collections import Counter
from imblearn.over_sampling import SMOTE
# summarize class distribution
print("Before oversampling: ", Counter(y_train))
```

Before oversampling: Counter({0: 604, 1: 86})

```
SMOTE = SMOTE()
# fit and apply the transform
X_train_SMOTE, y_train_SMOTE = SMOTE.fit_resample(X_train, y_train)
```

```
# summarize class distribution
print("After oversampling: ", Counter(y_train_SMOTE))
```

After oversampling: Counter({0: 604, 1: 604})

- We will split the training data into cross validation sets using “GridSearchCV” object, where number of folds is 10.

```
def grid_search_hyper(model, params, X_train, y_train, X_test):
    """This function for tuning a model and train it using GridSearchCV object
    inputs:
        model: the model you want to tune and train
        params: the hyperparameters that will be tuned
        X_train, y_train, X_test: data for train test
    outputs:
        grid.best_params_[1]: a dictionary of best hyperparameters for the model
        y_train_pred: the predicted train y
        y_test_pred: the predicted test y
        predicted_test_proba: the predicted probability test y
        grid: the trained GridSearchCV with best hyperparameters accessible"""

    grid = GridSearchCV(estimator=model,
                        param_grid=params,
                        scoring='accuracy',
                        cv=10,
                        return_train_score=False,
                        refit=True,
                        verbose=1)

    # training and Prediction
    grid.fit(X_train, y_train)
    # using the best estimator(the model with best hyperparameters) to predict the target using training and testing data
    y_train_pred = grid.best_estimator_.predict(X_train)
    y_test_pred = grid.best_estimator_.predict(X_test)
    # predict the probability of each sample to make smooth curve
    predicted_test_proba = grid.best_estimator_.predict_proba(X_test)

    return grid.best_params_, y_train_pred, y_test_pred, predicted_test_proba, grid
```

## Building Models and the Objective

- We train our models on a medical data that require a specific type of evaluation metrics such as Recall metric to be high on both two classes 0 (recovered) and 1 (death). So, our focus will be on this metric more than others to enhance it. The models we will use are:
  - (1) K-Nearest Neighbors (Without SMOTE)
  - (2) Logistic Regression (With SMOTE)
  - (3) Naïve Bayes (With SMOTE)
  - (4) Decision Trees (With SMOTE)
  - (5) Support Vector Machines (Without SMOTE)

- We noticed that when we use a data augmentation technique such as oversampling to make the dataset balanced with some algorithms the Recall metric becomes better at testing data.
- We will use SMOTE to oversample the data.

## 1. K-Nearest Neighbors (Without SMOTE)

- When we trained the model with GridSearchCV we got that the best hyperparameter is:
  - The best n\_neighbors is: 5
- **Recall**: in class 0 it's good, it's about 0.99, but in class 1 it's not good, it's about 0.64.
- **Precision**: in class 0 it's good, it's about 0.95, and in class 1, it's about 0.93.
- **F1-score**: in class 0 it's good, it's about 0.97, but in class 1 it's not good, it's about 0.76.
- **ROC/AUC**: it's good, it's about 0.99.


-Testing evaluation metrics:

```
1- The Recall is: 64%
1- The Precision is: 93%
1- The F1-score is: 76%
```

- The classification report for testing data

	precision	recall	f1-score	support
0	0.95	0.99	0.97	151
1	0.93	0.64	0.76	22
accuracy			0.95	173
macro avg	0.94	0.81	0.86	173
weighted avg	0.95	0.95	0.94	173

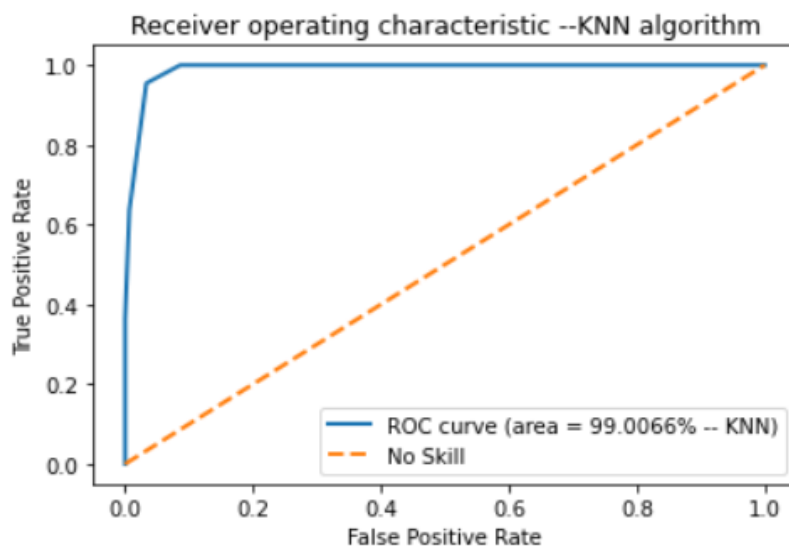
## The confusion matrix of Testing predicted data



```
: # plotting the ROC curve for KNN algorithm  
plot_roc(y_test,knn_test_proba_pred[:,1], title="KNN algorithm", algorithm="KNN")
```

AUC: 0.9901

No Skill AUC: 0.5000



- This model is doing well overall, but the Healthcare application requires a high Recall metric in two classes, where this model is high in class 0, but low in class 1.



## 2. Logistic Regression (With SMOTE)

- When we trained the model with GridSearchCV we got that the best hyperparameters are:
  - The best C is: 1.0
  - The best penalty is: l2
  - The best solver is: lbfgs
- **Recall**: in class 0 it's good, it's about 0.93, and in class 1, it's about 1.0.
- **Precision**: in class 0 it's good, it's about 1.0, but in class 1 it's not good, it's about 0.67.
- **F1-score**: in class 0 it's good, it's about 0.96, but in class 1 it's not good, it's about 0.80.
- **ROC/AUC**: it's good, it's about 0.98.


-Testing evaluation metrics:

```
1- The Recall is: 100%
1- The Precision is: 67%
1- The F1-score is: 80%
```

-----  
- The classification report for testing data

	precision	recall	f1-score	support
0	1.00	0.93	0.96	151
1	0.67	1.00	0.80	22
accuracy			0.94	173
macro avg	0.83	0.96	0.88	173
weighted avg	0.96	0.94	0.94	173

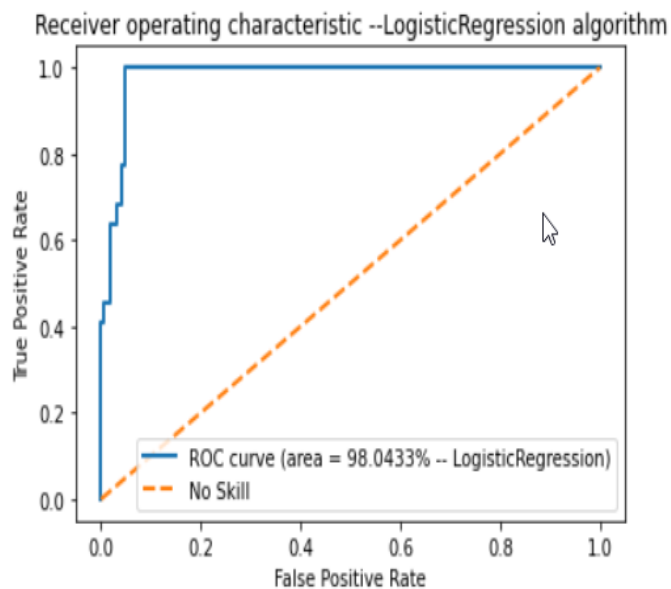
## The confusion matrix of Testing predicted data



```
# plotting the ROC curve for Logistic Regression algorithm  
plot_roc(y_test,logr_test_proba_pred[:,1], title="LogisticRegression algorithm", algorithm="LogisticRegression")
```

AUC: 0.9804

No Skill AUC: 0.5000



- This model is doing well overall, where this model is high in Recall metric with two classes.

### 3. Naïve Bayes (With SMOTE)

- We will use Gaussian Naïve Bayes.
- When we trained the model with GridSearchCV we got that the best hyperparameter is:
  - The best var\_smoothing is: 0.05542664520663104
- **Recall**: in class 0 it's good, it's about 0.88, and in class 1, it's about 1.0.
- **Precision**: in class 0 it's good, it's about 1.0, but in class 1 it's not good, it's about 0.55.
- **F1-score**: in class 0 it's good, it's about 0.94, but in class 1 it's not good, it's about 0.71.
- **ROC/AUC**: it's good, it's about 0.98.


-----  
-Testing evaluation metrics:

```
1- The Recall is: 100%
1- The Precision is: 55%
1- The F1-score is: 71%
```

-----  
- The classification report for testing data

	precision	recall	f1-score	support
0	1.00	0.88	0.94	151
1	0.55	1.00	0.71	22
accuracy			0.90	173
macro avg	0.78	0.94	0.82	173
weighted avg	0.94	0.90	0.91	173

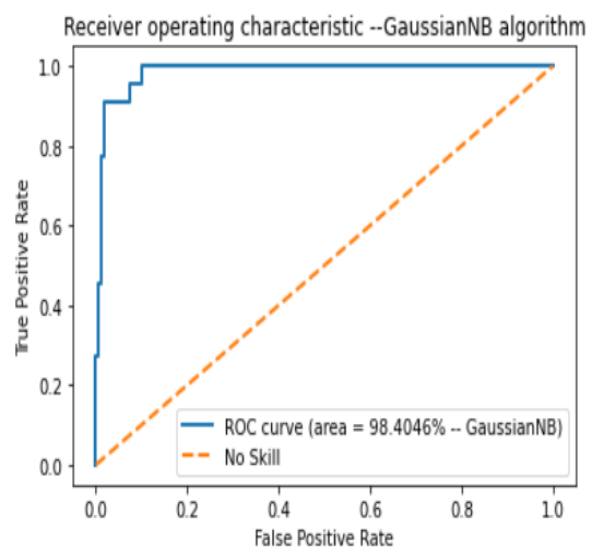
The confusion matrix of Testing predicted data



```
# plotting the ROC curve for GaussianNB algorithm  
plot_roc(y_test, gauNB_test_proba_pred[:,1], title="GaussianNB algorithm", algorithm="GaussianNB")
```

AUC: 0.9840

No Skill AUC: 0.5000



- This model is doing well overall, where this model is high in Recall metric with two classes.

## 2. Decision Trees (With SMOTE)

- When we trained the model with GridSearchCV we got that the best hyperparameters are:
  - The best criterion is: gini
  - The best max\_depth is: 8
  - The best min\_samples\_leaf is: 5
- **Recall**: in class 0 it's good, it's about 0.97, and in class 1, it's about 0.95.
- **Precision**: in class 0 it's good, it's about 0.99, and in class 1, it's about 0.81.
- **F1-score**: in class 0 it's good, it's about 0.98, and in class 1, it's about 0.88.
- **ROC/AUC**: it's good, it's about 0.96.


-Testing evaluation metrics:

```
1- The Recall is: 95%
1- The Precision is: 81%
1- The F1-score is: 88%
```

-----  
- The classification report for testing data

	precision	recall	f1-score	support
0	0.99	0.97	0.98	151
1	0.81	0.95	0.88	22
accuracy			0.97	173
macro avg	0.90	0.96	0.93	173
weighted avg	0.97	0.97	0.97	173

The confusion matrix of Testing predicted data

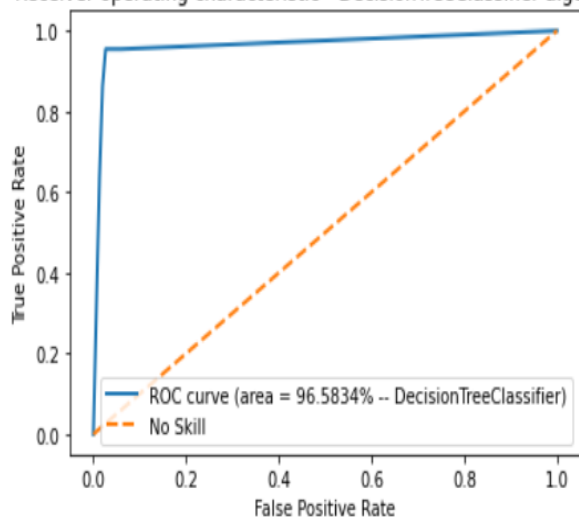


```
# plotting the ROC curve for DecisionTreeClassifier algorithm
plot_roc(y_test,decT_test_proba_pred[:,1], title="DecisionTreeClassifier algorithm", algorithm="DecisionTreeClassifier")
```

AUC: 0.9658

No Skill AUC: 0.5000

Receiver operating characteristic --DecisionTreeClassifier algorithm



- This model is doing well overall, where this model is high in Recall metric with two classes.

### 3. Support Vector Machines (Without SMOTE)

- When we trained the model with GridSearchCV we got that the best hyperparameters are:
  - The best C is: 1000
  - The best gamma is: 0.0001
- **Recall:** in class 0 it's good, it's about 1.0, and in class 1, it's about 0.91.
- **Precision:** in class 0 it's good, it's about 0.99, and in class 1, it's about 1.0.
- **F1-score:** in class 0 it's good, it's about 0.99, and in class 1, it's about 0.95.
- **ROC/AUC:** it's good, it's about 1.0.


-Testing evaluation metrics:

```
1- The Recall is: 91%
1- The Precision is: 100%
1- The F1-score is: 95%
```

-----  
- The classification report for testing data

	precision	recall	f1-score	support
0	0.99	1.00	0.99	151
1	1.00	0.91	0.95	22
accuracy			0.99	173
macro avg	0.99	0.95	0.97	173
weighted avg	0.99	0.99	0.99	173

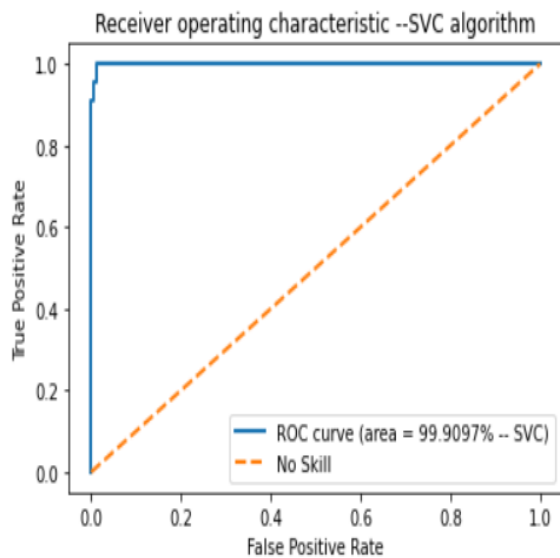
The confusion matrix of Testing predicted data



```
# plotting the ROC curve for SVC algorithm  
plot_roc(y_test,svc_test_proba_pred[:,1], title="SVC algorithm", algorithm="svc")
```

AUC: 0.9991

No Skill AUC: 0.5000





- This model is doing well overall, where this model is high in Recall metric with two classes.

## Conclusion

All models are doing well overall, but in healthcare applications that require a high recall metric on the class 1 most of the time, and in our case class 1 (death class), so, the best model should have a high Recall value on this class, and be balanced with other metrics, as shown above Support Vector Machine is suitable for this.

- Best model with each evaluation metric:
  - Recall: Logistic Regression, Decision Tree, and Support vector Machines.
  - Precision: K-nearest Neighbor, and Support Vector Machines.
  - F1-score: Support Vector Machines.
  - ROC/AUC: Support Vector Machines.

Receiver operating characteristic --

